

# Project 3 Discussion

Multi-Process and Multi-Threaded Web Servers

# Setup

- Use the same virtual machine as in Project 2
  - Login as *root*
- If you want a new copy, grab the appropriate software and disk image from the assignment page

# Setup

As the root user:

```
apt-get update  
apt-get install tmux links
```

To get the single-threader server code from thoth:

```
scp USERNAME@thoth.cs.pitt.edu:/u/OSLab/original/server.c .
```

Now, let's try to send requests to  
the single-threaded server  
(server.c)

# Before that, do this

Just type *tmux* on the command line and press ENTER

1. Run `tmux` ←
2. Execute your server
3. Hit **CTRL+B**, release, and then type " (a quotation mark -- this requires holding shift down to type)
4. You should see a window with a shell appear at the bottom
5. At any point, you can switch which shell has focus by typing **CTRL+B** and then pressing an arrow key to move up or down amongst the windows.
6. When you're done testing, type `exit` to leave the shell and the bottom window will go away. You'll probably want to use **CTRL+C** to kill your server. Then you can type `exit` again to leave `tmux`.

```
root@debianvm:~#
```

terminal 1

```
root@debianvm:~#
```

terminal 2

```
[0] 0: bash*
```

```
"debianvm" 23:22 19-Oct-20
```

# Running server and client

- Run the server on one terminal
- Run the client(s) on the other terminal

DEMO



```
root@debianvm:~# gcc -o server server.c
root@debianvm:~# ./server
GET /page3.html HTTP/1.1
User-Agent: Wget/1.21
Accept: */*
Accept-Encoding: identity
Host: 127.0.0.1
Connection: Keep-Alive
```

```
root@debianvm:~# wget http://127.0.0.1/page3.html
--2023-10-19 23:27:32-- http://127.0.0.1/page3.html
Connecting to 127.0.0.1:80... connected.
HTTP request sent, awaiting response... 404 Not Found
2023-10-19 23:27:32 ERROR 404: Not Found.
```

```
root@debianvm:~#
```

```
[0] 0:bash*
```

```
"debianvm" 23:27 19-Oct-23
```

- Server runs on this terminal.
  - Server could not find the requested file, page3.html
- 
- Client runs on this terminal.
  - Client requested for page3.html
    - Server returned status 404(file not found).

# What do you need to do?

- Write code for a **multi-process** version of the server in the file *server\_proc.c*
- Write code for a **multi-threaded** version of the server in the file *server\_thread.c*
- Write a **1-2 page paper**
  - Answer this question
    - Which implementation of the server is faster?
  - Describe your experimental design, your results, and your conclusions in a 1-2 page paper.

Multi-process version of the  
server

# *server\_proc.c*

- Use *fork()* to create a **worker process**
- Worker Process
  - Handles **sending** the requested file over the network connection
  - When done sending the file, it should log the request for that particular webpage by appending it to a file named *stats\_proc.txt*
    - Access to *stats\_proc.txt* needs to be synchronized
      - Use *sem\_wait()* and *sem\_post()* from <semaphore.h>

exclusively accessed, so you'll need to do some sort of synchronization. Log the **name** of the file requested, its **size in bytes**, the **ending CPU time** of the request and the **elapsed CPU time** in seconds of the request, in the following tab-separated format:

```
file1.html    13452  0.5123  0.0123
```

## *server\_proc.c*

- If you need shared space to hold something between processes, use **mmap()** with the MAP\_SHARED and MAP\_ANONYMOUS flags.
- When you create a child process, you'll have two copies of the **socket file descriptor** and two of the **connection file descriptors**. The parent only needs one of those and the child only needs one. Make sure to close() the other.

# Using fork() to create a child process

- DEMO

# Using fork() to create a child process

```
#include<stdio.h>
#include <stdlib.h>
#include <unistd.h>
int main()
{
    int x = 10;

    if(fork() == 0)
    {
        printf("[Child] pid %d from [parent] pid %d\n",getpid(),getppid());

        x = 20;
        printf("Child x : %d\n", x);

        exit(0);
    }

    wait(NULL);
    printf("Parent x: %d\n", x);
}
```

Multi-threaded version of the  
server



# *server\_thread.c*

- Use *pthread\_create()* to create a **worker thread**
- Worker Thread
  - Handles **sending** the requested file over the network connection
  - When done sending the file, it should log the request for that particular webpage by appending it to a file named *stats\_thread.txt*
    - Access to *stats\_thread.txt* needs to be synchronized
      - Use *pthread\_mutex\_lock()* and *pthread\_mutex\_unlock()* from *<pthread.h>*

exclusively accessed, so you'll need to do some sort of synchronization. Log the **name** of the file requested, its **size** in bytes, the **ending CPU time** of the request and the **elapsed CPU time** in seconds of the request, in the following **tab-separated format**:

```
file1.html    13452  0.5123  0.0123
```

## *server\_thread.c*

- Sharing between threads is just allocating things in the heap or globals.
- Compile your `server_thread.c` file with the `-pthread` flag
  - DEMO

CLIENT

# Test from the client side

- `nc localhost 80`
- `wget http://127.0.0.1/page.html`
- `links http://127.0.0.1/page.html`

# Client

- You can run **multiple programs** from the client side to send requests to the server in parallel
  - To run two programs simultaneously on the command line, put an & between them
    - Here, a program can be wget or curl or any program that can download from a web server. Or you can even write your own test program as HTTP is not a complicated protocol.