

# Project 4 Discussion

DIY Semaphores

# Setup

- Same as Project 2
- So, start early!!!!

## **Setting up, building, and installing the Kernel**

Follow the exact same steps as in project 2. I'd suggest starting from the original kernel source (you can download/extract [if you kept the download] a new VM image if you want, or simply delete the old linux kernel folder and extract the source anew).

Note this means you'll have the same two hour-long builds as anytime we add system calls the entire kernel will be rebuilt. Make sure you start this setup early.

# Overall

- Simulate producer-consumer problem
- Write the user space code in :
  - prodcons.c
- Add new syscalls in the kernel space by modifying the following:
  - sys.c
  - syscalls.h
  - unistd.h
  - syscall\_64.tbl
  - syscall\_32.tbl

sys.C

# Semaphore

```
struct cs1550_sem
{
    int value;
    struct mutex *lock;    //So the kernel can lock on this instance
    //Some process queue of your devising
};
```

# What is *struct mutex*?

```
struct cs1550_sem
{
    int value;
    struct mutex *lock;    //So the kernel can lock on this instance
    //Some process queue of your devising
};
```

- Predefined struct defined in *linux/mutex.h*
- *lock* will be used for maintaining atomicity in your up() and down() syscalls

# What should be the process queue?

- A queue for storing processes
- What type of data should each node of the queue store?
  - *struct task\_struct* is the structure for process metadata
  - The global variable *current* is a **pointer** to the current process
    - Each node should be able to store *current*
    - So, the data structure of each node in the queue should be?
- You can build a queue of your own devising
  - Hint: You can use Linked List with head and tail

# Syscalls



```
asmlinkage long sys_cs1550_seminit(struct  
cs1550_sem *sem, int value)
```

- Initialize the semaphore value
- Initialize `sem->lock`
  - Allocate memory for `sem->lock`
    - Use `kmalloc`
  - Initialize using `mutex_init(sem->lock)`

asmlinkage long **sys\_cs1550\_down**(struct  
cs1550\_sem \*sem)

```
void down () {  
    value -= 1;  
    if (value < 0) {  
        // add this process to pl  
        pl.enqueue(currentProcess);  
        Sleep();  
    }  
}
```

Enqueue current  
process (maintained  
by global pointer  
variable *current*) to  
the tail of your queue

asmlinkage long sys\_cs1550\_down(struct  
cs1550\_sem \*sem)

```
void down () {  
    value -= 1;  
    if (value < 0) {  
        // add this process to pl  
        pl.enqueue(currentProcess);  
        Sleep();  
    }  
}
```

Make the current process go to sleep

- set\_current\_state(TASK\_INTERRUPTIBLE);
- schedule();

```
asmlinkage long sys_cs1550_down(struct  
cs1550_sem *sem)
```

```
    sys_cs1550_down(struct cs1550_sem *sem)  
{  
    mutex_lock(sem->lock)
```

```
    //CRITICAL REGION
```

```
    mutex_unlock(sem->lock)  
}
```

Please note that you may need to call mutex\_unlock() more than once if there is some part of your code that is not reachable

# Do not return without unlocking (Example)

```
if {  
    ...  
    return x;  
}  
mutex_unlock(sem->lock);  
return y;
```



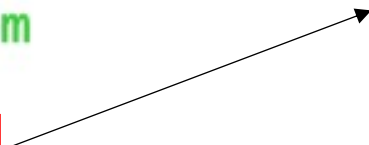
```
result = y;  
if {  
    ...  
    result = x;  
}  
mutex_unlock(sem->lock);  
  
return result;
```



asmlinkage long sys\_cs1550\_up(struct  
cs1550\_sem \*sem)

```
void up () {  
    Process P;  
    value += 1;  
    if (value <= 0) {  
        // remove a process P from  
        pl  
        P = pl.dequeue();  
        Wakeup(P);  
    }  
}
```

Dequeue the process  
from the head of the  
queue



asmlinkage long **sys\_cs1550\_up**(struct  
cs1550\_sem \*sem)

```
void up () {  
    Process P;  
    value += 1;  
    if (value <= 0) {  
        // remove a process P from  
        pl  
        P = pl.dequeue();  
        Wakeup(P);  
    }  
}
```

Use

*wake\_up\_process(sleeping\_task);*

- *sleeping\_task* is the process that you just dequeued
- Type *struct task\_struct\**

```
asmlinkage long sys_cs1550_up(struct  
cs1550_sem *sem)
```

```
    sys_cs1550_up(struct cs1550_sem *sem)  
{  
    mutex_lock(sem->lock)
```

```
    //CRITICAL REGION
```

```
    mutex_unlock(sem->lock)  
}
```

Please note that you may need to call `mutex_unlock()` more than once if there is some part of your code that is not reachable



# Do not return without unlocking (Example)

```
if {  
    ...  
    return x;  
}  
mutex_unlock(sem->lock);  
return y;
```



```
result = y;  
if {  
    ...  
    result = x;  
}  
mutex_unlock(sem->lock);  
  
return result;
```



prodcons.c

# mmap()

- call mmap() once to store
  - Shared semaphore variables
  - Shared buffer
- Use MAP\_SHARED|MAP\_ANONYMOUS
- Size = Size of three semaphore variables + Size of buffer
- Use pointer arithmetic to access the mmap'd space
  - *struct cs1550\_sem\* empty = ptr*
  - *struct cs1550\_sem\* full = empty + 1*
  - and so on ...

# prodcons.c – producer and consumer

```
void *producer(void *junk) {  
    while(1) {  
        sem_wait(&semempty);  
        sem_wait(&semmutex);  
  
        buffer[in] = total++;  
        printf("Produced: %d\n", buffer[in]);  
        in = (in + 1) % N;  
        counter++;  
  
        sem_post(&semmutex);  
        sem_post(&semfull);  
    }  
}
```

```
void *consumer(void *junk) {  
    while(1) {  
        sem_wait(&semfull);  
        sem_wait(&semmutex);  
  
        printf("Consumed: %d\n", buffer[out]);  
        out = (out + 1) % N;  
        counter--;  
  
        sem_post(&semmutex);  
        sem_post(&semempty);  
    }  
}
```

You can consider this as kind of a Pseudocode for your implementation. Make sure to use the up() and down() syscalls that you added in sys.c

# Important Points

- For producer
  - `fork()` as many times as the number of producers
  - Each child process should execute the producer part of your code
- For consumer
  - `fork()` as many times as the number of consumers
  - Each child process should execute the consumer part of your code
- Call the `wait()` function in the parent process to wait for the child processes to finish execution